

YUXUAN ZHANG

858-729-3152 | yuxuan_zhang13@outlook.com | linkedin.com/in/yuxuan-zhang-dexter | github.com/Yuxuan-Zhang-Dexter

EDUCATION

University of California, San Diego

Sep. 2025 – Expected Jun. 2027

Master of Science in Computer Science

University of California, San Diego

Sep. 2021 – March 2025

Bachelor of Science in Data Science

GPA: 4.0 / 4.0, **Summa Cum Laude** (Top 2%, highest honor)

SELECTED HIGHLIGHTS

• Publications

- **lmgame-Bench: How Good are LLMs at Playing Games?** (NeurIPS '25 — Submitted) Lanxiang Hu, Mingjia Huo, [Yuxuan Zhang](#), Haoyang Yu, Eric P. Xing, Ion Stoica, Tajana Rosing, Haojian Jin, Hao Zhang
- **General Modular Harness for LLM Agents in Multi-Turn Gaming Environments** (ICML '25 Workshop) [Yuxuan Zhang](#), Haoyang Yu, Lanxiang Hu, Haojian Jin, Hao Zhang

• Open-Source Projects & Media

- **AI GameAgent** — Co-led an open-source AI-agent framework for classic games (718★)
github.com/lmgame-org/GamingAgent
- **Agentic Multi-Turn RL Training Library** — Co-designed a scalable multi-turn RL library for VLMs/LLMs, enabling synchronous rollouts and agent-centric training pipelines
github.com/mignonjia/lmgame_train_reorg
- **AI Ace Attorney Gameplay** — Led AI-agent design and demo that reached 950 k views on X/Twitter
x.com/haoailab/status/1912231343372812508

PROFESSIONAL EXPERIENCE

HDSI Research Assistant

Mar. 2025 – Present

University of California, San Diego, HDSI Department, Dr. Hao Zhang

La Jolla, CA

- Designed and Deployed Gradio App on Hugging Face to visualize performance of 16 AI models across 7 games.
- Led modularization of AI agent architecture for games including 2048, Sokoban, Tetris, Super Mario Bros...
- Building an end-to-end gameplay and evaluation pipeline using Gymnasium (in progress).

Undergraduate Research Assistant, Hao AI Lab

Jun. 2024 – Mar. 2025

University of California, San Diego, HDSI Department, Dr. Hao Zhang

La Jolla, CA

- Designed and deployed the *Game Arena* website (lmgame.org) to benchmark LLMs with real-player gameplay.
- Integrated PostgreSQL to compute and update ELO scores for models and users, optimized for scale.
- Developed a multithreaded MySQL + AWS S3 pipeline to process 6.9 TB of logs, achieving a 4× throughput gain.

Data Scientist Research Intern, Maximov Lab

Apr 2023 – Dec 2024

Scripps Research Institute, Neuroscience Department, Dr. Anton Maximov

La Jolla, CA

- Built a 3D EM neuron segmentation pipeline with Florence/SAM2 models, achieving $\leq 80\%$ validation accuracy.
- Developed an interactive web app using NeuronGlancer for 3D visualization of 32-bit large-scale neuron datasets.
- Converted 500+ GB of EM images into compressed HDF5, cutting storage by 70%.

TECHNICAL SKILLS

Programming Languages: Python, Java, C/C++, SQL (Postgres), JavaScript, HTML/CSS, R, Lua

Frameworks & Libraries: Pytorch, Ray, DeepSpeed, Gymnasium, Triton, Node.js, PIL, cv2, React, OpenCL

Developer Tools: Git, Docker, Kubernetes, PostgreSQL, MySQL, AWS S3, Wandb, MongoDB, Postman, Linux/UI

AI API Integration: OpenAI, Google Gemini, Anthropic, Grok, Deepseek, Together AI, and LLaMA

HONORS & AWARDS

Cornell, Maryland, Max Planck Predoctoral Research School

Aug. 2023

Joined a week-long research talk on software systems and informatics at the Max Planck Institute, Germany.

UCSD HDSI Scholarship Program

Sep. 2023 – Jun. 2024

Awarded a scholarship for a research project proposal, receiving financial support for the academic year.